

Lectures 8: Class demo

Contents

- Imports
- Demo: Model interpretation of linear classifiers
- ? ? Questions for you

Imports

```
import os
import sys

sys.path.append(os.path.join(os.path.abspath(".."), (".."), "code"))
from utils import *
import matplotlib.pyplot as plt
import mglearn
import numpy as np
import pandas as pd
from plotting_functions import *
from sklearn.dummy import DummyClassifier
from sklearn.feature_extraction.text import CountVectorizer, TfidfVectorizer
from sklearn.model_selection import cross_val_score, cross_validate, train_test_split
from sklearn.pipeline import make_pipeline
from sklearn.linear_model import LogisticRegression
%matplotlib inline
pd.set_option("display.max_colwidth", 200)
DATA_DIR = os.path.join(os.path.abspath(".."), (".."), "data/")
```

Demo: Model interpretation of linear classifiers

- One of the primary advantage of linear classifiers is their ability to interpret models.
- For example, with the sign and magnitude of learned coefficients we could answer questions such as which features are driving the prediction to which direction.

- We'll demonstrate this by training `LogisticRegression` on the famous [IMDB movie review](#) dataset. The dataset is a bit large for demonstration purposes. So I am going to put a big portion of it in the test split to speed things up.

```
imdb_df = pd.read_csv(DATA_DIR + "imdb_master.csv", encoding="ISO-8859-1")
imdb_df.head()
```

	review	sentiment
0	One of the other reviewers has mentioned that after watching just 1 Oz episode you'll be hooked. They are right, as this is exactly what happened with me. The first thing that struck me...	positive
1	A wonderful little production. The filming technique is very unassuming- very old-time-BBC fashion and gives a comforting, and sometimes discomforting, sense of realism to the entire p...	positive
2	I thought this was a wonderful way to spend time on a too hot summer weekend, sitting in the air conditioned theater and watching a light-hearted comedy. The plot is simplistic, but the dialogue i...	positive
3	Basically there's a family where a little boy (Jake) thinks there's a zombie in his closet & his parents are fighting all the time. This movie is slower than a soap opera... and suddenl...	negative
4	Petter Mattei's "Love in the Time of Money" is a visually stunning film to watch. Mr. Mattei offers us a vivid portrait about human relations. This is a movie that seems to be telling us what mone...	positive

Let's clean up the data a bit.

```
import re

def replace_tags(doc):
    doc = doc.replace(r"<br />", " ")
    doc = re.sub(r"https://\S*", "", doc)
    return doc
```

```
imdb_df["review_pp"] = imdb_df["review"].apply(replace_tags)
```

Are we breaking the Golden rule here?

Let's split the data and create bag of words representation.

```
train_df, test_df = train_test_split(imdb_df, test_size=0.9, random_state=123)
X_train, y_train = train_df["review_pp"], train_df["sentiment"]
X_test, y_test = test_df["review_pp"], test_df["sentiment"]
train_df.shape
```

```
(5000, 3)
```

```
vec = CountVectorizer(stop_words="english")
bow = vec.fit_transform(X_train)
bow
```

```
<Compressed Sparse Row sparse matrix of dtype 'int64'
with 439384 stored elements and shape (5000, 38867)>
```

Examining the vocabulary

- The vocabulary (mapping from feature indices to actual words) can be obtained using `get_feature_names_out()` on the `CountVectorizer` object.

```
vocab = vec.get_feature_names_out()
```

```
vocab.shape
```

```
(38867,)
```

```
vocab[0:10] # first few words
```

```
array(['00', '000', '007', '0079', '0080', '0083', '00pm', '00s', '01',
      '0126'], dtype=object)
```

```
vocab[2000:2010] # some middle words
```

```
array(['apprehensive', 'apprentice', 'approach', 'approached',
      'approaches', 'approaching', 'appropriate', 'appropriated',
      'appropriately', 'approval'], dtype=object)
```

```
vocab[:,500] # words with a step of 500
```

```
array(['00', 'aaja', 'affection', 'ambrosine', 'apprehensive', 'attract',
      'barbara', 'bereavement', 'blore', 'brazenly', 'businessman',
      'carrel', 'chatterjee', 'claudio', 'commanding', 'consumed',
      'cramped', 'cynic', 'defining', 'deviates', 'displaced',
      'dramatized', 'edie', 'enforced', 'evolving', 'fanatically',
      'fingertips', 'formal', 'gaffers', 'giogio', 'gravitas',
      'halliday', 'heist', 'hoot', 'iliad', 'infiltrate', 'investment',
      'jobson', 'kidnappee', 'landsbury', 'licentious', 'lousiest',
      'malã', 'maã', 'mice', 'molla', 'museum', 'newtonian',
      'obsessiveness', 'outbursts', 'parapsychologist', 'perpetuates',
      'plasters', 'powers', 'property', 'rabies', 'reclined', 'renters',
      'ridiculous', 'rube', 'sayid', 'select', 'shivers', 'skinheads',
      'sohail', 'spot', 'stomaches', 'suitcase', 'syrupy', 'terrorist',
      'tolerance', 'triangular', 'unbidden', 'unrevealed', 'verneuil',
      'walrus', 'wilcox', 'xxxxviii'], dtype=object)
```

```
y_train.value_counts()
```

```
sentiment
positive    2517
negative    2483
Name: count, dtype: int64
```

Model building on the dataset

First let's try `DummyClassifier` on the dataset.

```
dummy = DummyClassifier()
scores = cross_validate(dummy, X_train, y_train, return_train_score=True)
pd.DataFrame(scores)
```

	fit_time	score_time	test_score	train_score
0	0.001099	0.000938	0.504	0.50325
1	0.000945	0.000799	0.504	0.50325
2	0.000909	0.000825	0.503	0.50350
3	0.000986	0.000831	0.503	0.50350
4	0.000915	0.000827	0.503	0.50350

We have a balanced dataset. So the `DummyClassifier` score is around 0.5.

Now let's try logistic regression.


```
pipe_lr = make_pipeline(
    CountVectorizer(stop_words="english"),
    LogisticRegression(max_iter=1000),
)
scores = cross_validate(pipe_lr, X_train, y_train, return_train_score=True)
pd.DataFrame(scores)
```

	fit_time	score_time	test_score	train_score
0	0.385554	0.059766	0.828	0.99975
1	0.377250	0.061282	0.830	0.99975
2	0.383834	0.059718	0.848	0.99975
3	0.370757	0.058832	0.833	1.00000
4	0.372611	0.062104	0.840	0.99975

Seems like we are overfitting. Let's optimize the hyperparameter `C` of LR and `max_features` of `CountVectorizer`.

```
from scipy.stats import loguniform, randint, uniform
from sklearn.model_selection import RandomizedSearchCV

param_dist = {
    "countvectorizer__max_features": randint(10, len(vocab)),
    "logisticregression__C": loguniform(1e-3, 1e3)
}
pipe_lr = make_pipeline(CountVectorizer(stop_words="english"), LogisticRegression)
random_search = RandomizedSearchCV(pipe_lr, param_dist, n_iter=10, n_jobs=-1,
random_search.fit(X_train, y_train)
```



```
pd.DataFrame(random_search.cv_results_)
```

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_count
0	0.588244	0.021243	0.110041	0.006330	
1	0.657085	0.028550	0.110500	0.006087	
2	0.501113	0.022524	0.105766	0.008775	
3	0.519245	0.042000	0.126679	0.018991	
4	0.503987	0.069851	0.122257	0.018959	
5	0.552576	0.008925	0.106016	0.012166	
6	0.732404	0.057531	0.130433	0.014476	
7	0.743628	0.077521	0.100707	0.004814	
8	0.572081	0.014627	0.104338	0.004348	
9	0.468955	0.052592	0.088705	0.014833	

10 rows × 22 columns

```
cv_scores = random_search.cv_results_['mean_test_score']
train_scores = random_search.cv_results_['mean_train_score']
countvec_max_features = random_search.cv_results_['param_countvectorizer__max_
```


```
pd.DataFrame(random_search.cv_results_)[
    [
        "mean_test_score",
        "mean_train_score",
        "param_logisticregression__C",
        "param_countvectorizer__max_features",
        "mean_fit_time",
        "rank_test_score",
    ]
].set_index("rank_test_score").sort_index()
```

	mean_test_score	mean_train_score	param_logisticregression__
rank_test_score			
1	0.8400	0.99700	0.17835
2	0.8396	0.99925	0.31768
3	0.8386	0.96020	0.02784
4	0.8346	1.00000	356.0177
5	0.8318	1.00000	5.46987
6	0.8246	1.00000	6.81927
7	0.8238	0.88735	0.00470
8	0.8236	1.00000	6.84976
9	0.8178	0.87315	0.00292
10	0.8004	1.00000	12.79360

Let's train a model on the full training set with the optimized hyperparameter values.

```
best_model = random_search.best_estimator_
```

```
best_model
```

Examining learned coefficients

- The learned coefficients are exposed by the `coef_` attribute of [LogisticRegression](#) object.

```
# Get feature names
feature_names = best_model.named_steps['countvectorizer'].get_feature_names_out()

# Get coefficients
coeffs = best_model.named_steps["logisticregression"].coef_.flatten()
```

```
word_coeff_df = pd.DataFrame(coeffs, index=feature_names, columns=["Coefficient"])
word_coeff_df
```

	Coefficient
00	0.067353
000	0.105312
007	0.006590
0083	0.027134
00s	-0.052512
...	...
zwick	0.019599
zyada	-0.026069
zzzzip	-0.000101
zzzzz	-0.031992
â½	-0.010470

30575 rows × 1 columns

- Let's sort the coefficients in descending order.
- Interpretation
 - if $w_j > 0$ then increasing x_{ij} moves us toward predicting $+1$.
 - if $w_j < 0$ then increasing x_{ij} moves us toward predicting -1 .

```
word_coeff_df.sort_values(by="Coefficient", ascending=False)
```


	Coefficient
excellent	0.876824
perfect	0.713843
amazing	0.707115
wonderful	0.685750
surprised	0.679117
...	...
waste	-0.823069
boring	-0.847032
terrible	-0.847147
awful	-1.026551
worst	-1.303253

30575 rows × 1 columns

- The coefficients make sense!

Let's visualize the top 20 features.

```
mglearn.tools.visualize_coefficients(coeffs, feature_names, n_top_features=20)
```



Let's explore prediction of the following new review.

```
fake_reviews = ["It got a bit boring at times but the direction was excellent  

    "The plot was shallower than a kiddie pool in a drought, but hey, at least we  

    "]
```

Let's get prediction probability scores of the fake review.

```
best_model.predict(fake_reviews)
```

```
array(['positive', 'negative'], dtype=object)
```

```
# Get prediction probabilities for fake reviews  
best_model.predict_proba(fake_reviews)
```

```
array([[0.13615126, 0.86384874],  

       [0.72517628, 0.27482372]])
```

```
best_model.classes_
```

```
array(['negative', 'positive'], dtype=object)
```

We can find which of the vocabulary words are present in this review:


```
def plot_coeff_example(model, review, coeffs, feature_names, n_top_feats=6):
    print(review)
    feat_vec = model.named_steps["countvectorizer"].transform([review])
    words_in_ex = feat_vec.toarray().ravel().astype(bool)

    ex_df = pd.DataFrame(
        data=coeffs[words_in_ex],
        index=np.array(feature_names)[words_in_ex],
        columns=["Coefficient"],
    )
    mglearn.tools.visualize_coefficients(
        coeffs[words_in_ex], np.array(feature_names)[words_in_ex], n_top_features=
    )
    return ex_df.sort_values(by=["Coefficient"], ascending=False)
```

```
plot_coeff_example(best_model, fake_reviews[0], coeffs, feature_names)
```

It got a bit boring at times but the direction was excellent and the acting was


	Coefficient
excellent	0.876824
overall	0.492946
enjoyed	0.481338
bit	0.397143
highly	0.386792
times	0.287180
recommend	0.182552
flawless	0.071724
movie	-0.035484
got	-0.097060
direction	-0.146983
acting	-0.213584
boring	-0.847032



```
plot_coeff_example(best_model, fake_reviews[1], coeffs, feature_names)
```

The plot was shallower than a kiddie pool in a drought, but hey, at least we no

	Coefficient
know	0.256913
stick	0.029108
kiddie	-0.062415
big	-0.077145
screen	-0.078166
hey	-0.083086
pool	-0.123016
plot	-0.276647
avoid	-0.567104



Most positive review

- Remember that you can look at the probabilities (confidence) of the classifier's prediction using the `model.predict_proba` method.
- Can we find the reviews where our classifier is most certain or least certain?

```
# only get probabilities associated with pos class
pos_probs = best_model.predict_proba(X_train)[
    :, 1
] # only get probabilities associated with pos class
pos_probs
```

```
array([0.98488155, 0.17244179, 0.96027595, ..., 0.80965294, 0.91813092,
       0.00243243])
```

What's the index of the example where the classifier is most certain (highest `predict_proba` score for positive)?

```
most_positive_id = np.argmax(pos_probs)
```

```
print("True target: %s\n" % (y_train.iloc[most_positive_id]))
print("Predicted target: %s\n" % (best_model.predict(X_train.iloc[[most_positi
print("Prediction probability: %0.4f" % (pos_probs[most_positive_id]))
```

True target: positive

Predicted target: positive

Prediction probability: 1.0000

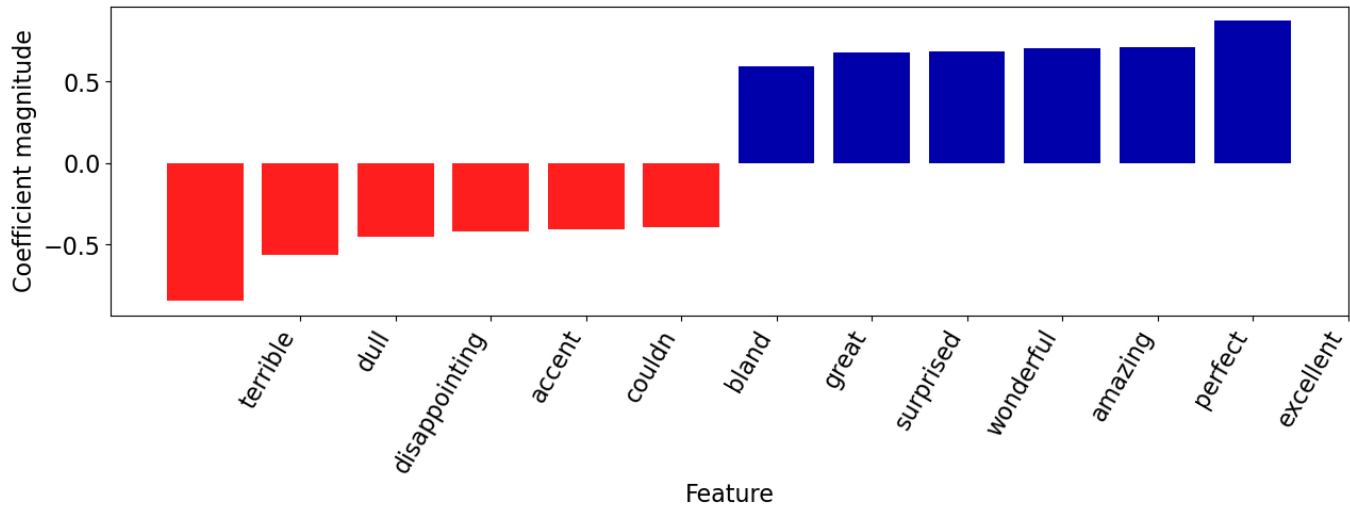
Let's examine the features associated with the review.

```
plot_coeff_example(best_model, X_train.iloc[most_positive_id], coeffs, feature
```

This is an awesome Amicus horror anthology, with 3 great stories, and fantastic

	Coefficient
excellent	0.876824
perfect	0.713843
amazing	0.707115
wonderful	0.685750
surprised	0.679117
...	...
couldn	-0.409726
accent	-0.419552
disappointing	-0.451959
dull	-0.565241
terrible	-0.847147

173 rows x 1 columns



The review has both positive and negative words but the words with **positive** coefficients win in this case!

Most negative review

```
neg_probs = best_model.predict_proba(X_train)[
    :, 0
] # only get probabilities associated with neg class
neg_probs
```

```
array([0.01511845, 0.82755821, 0.03972405, ..., 0.19034706, 0.08186908,
       0.99756757])
```

```
most_negative_id = np.argmax(neg_probs)
```

```
print("Review: %s\n" % (X_train.iloc[[most_negative_id]]))
print("True target: %s\n" % (y_train.iloc[most_negative_id]))
print("Predicted target: %s\n" % (best_model.predict(X_train.iloc[[most_negative_id]])))
print("Prediction probability: %0.4f" % (neg_probs[most_negative_id]))
```

```
Review: 13452    Zombi 3 starts as a group of heavily armed men steal a experim
Name: review_pp, dtype: object
```

```
True target: negative
```

```
Predicted target: negative
```


```
Prediction probability: 1.0000
```

```
plot_coeff_example(best_model, X_train.iloc[most_negative_id], coeffs, feature
```

```
Zombi 3 starts as a group of heavily armed men steal a experimental chemical de
```


	Coefficient
fun	0.434700
makes	0.407873
entertaining	0.407167
late	0.299810
right	0.272165
...	...
bad	-0.652233
poor	-0.712592
boring	-0.847032
terrible	-0.847147
awful	-1.026551

317 rows × 1 columns



The review has both positive and negative words but the words with negative coefficients win in this case!

? ? Questions for you

Question for you to ponder on

- Is it possible to identify most important features using k -NNs? What about decision trees?