

Lecture 1 - Frequentist and Bayesian Overview, Probabilistic Generative Models, and Stan

Contents

- High-Level Goals of this Course
- Course Overview
- Textbook
- Today's Learning Objectives
- Loading Libraries
- 1. Review of Frequentist Statistical Inference
- 2. Bayesian Statistics
- 3. Generative Models
- 4. **Stan** and **rstan** Basics
- 5. Difference between Probability and Likelihood
- 6. Wrapping Up with Frequentist Drawbacks

High-Level Goals of this Course

- Use Bayesian reasoning when modeling data.
- Apply Bayesian statistics to regression models.
- Compare and contrast Bayesian and frequentist methods, and evaluate their relative strengths.
- Use appropriate statistical libraries and packages for performing Bayesian inference.

Course Overview

- **Non-project MDS course:** eight lectures and four labs.
- Our focus is **model-building, computation, and interpretation of results.**
- We will build models in `Stan` along with `rstan`.
- Knowing how and when to use different statistical distributions is a **great asset.**
- `R` + `Stan` for lectures and labs.

Textbook

We will use a textbook in this course. Its name is [Bayes Rules! An Introduction to Applied Bayesian Modeling](#)). We will be posting suggested and optional readings of this book before our lecture time.

Today's Learning Objectives

1. Review statistical inference (frequentist so far!).
2. Pave the way to Bayesian statistics.
3. Introduce probabilistic generative models.
4. Illustrate the basic use of `Stan` and `rstan` via Monte Carlo simulations.
5. Differentiate probability and likelihood in Statistics.

Loading Libraries

```
library(tidyverse)
library(infer)
library(tidyverse)
library(cowplot)
library(datateachr)
library(bayesrules)
```

1. Review of Frequentist Statistical Inference

In the frequentist courses, we have used **observed data** (coming from a random sample) to **estimate** and **characterize uncertainty** in **unknown** (or **latent**) **population** quantities of interest. For instance:

- An **unknown** population mean μ .
- An **unknown** population variance σ^2 .
- An **unknown** population median M .

However, what do we mean when we say **latent population** quantities? Let us find it out.

1.1. Latent Variables

These **latent quantities** may be real **but not directly observable**. Hence, they are linked to other **observable variables**. For example:

- Using **online ad click data** to estimate **the total lifetime revenue**.
- Using **genome sequencing** to infer **the origin of a virus during an outbreak**.
- Using **robotic Light Detection and Ranging (LiDAR) sensors** to estimate **the robot's position**.

Or **completely hypothetical**:

- Using **tennis game win/loss data** to infer a **ranking of players**.
- Using **text data** to learn **the underlying hierarchy of topics**.



! Important

For these inquiries, the frequentist statistician's usual hammers are **point estimates** and **confidence intervals**.

1.2. Why Care About Uncertainty?


Cannot we just collect more data?

Not if it is **expensive** (e.g., rocket telemetry), **fundamentally limited** (e.g., robot odometry), or **outright impossible** (e.g., ancestral species).



1.3. What is a Confidence Interval (CI)?

If we care about uncertainty in a **frequentist** inferential approach, what is a confidence interval?



Let us start with the first **in-class question**.



Exercise 1

Answer TRUE or FALSE:

Assuming a random sample of size n composed of the random variables X_1, X_2, \dots, X_n (not observed sampled values!), a frequentist 90%-confidence interval (CI) is a random interval that contains an unknown fixed population parameter of interest with probability 90% before observing the random data.

A. TRUE

B. FALSE

Solution to [Exercise 1](#)

It is **TRUE**. The above statement **theoretically** defines the frequentist random CI in the general context of a random sample composed of n random variables X_1, X_2, \dots, X_n . Then in practice, once we **observe** the sampled data, the 90%-CI will either contain or not the unknown **fixed** population parameter of interest, which we aim to infer.

Then, suppose you have a large enough sampling budget to draw a considerable amount of m random samples (i.e., **we rely on a large frequency of events!**) to make inference on your population parameter of interest and compute their corresponding frequentist 90%-CIs. Under the frequentist paradigm, approximately 90% of these m CIs will contain the true **fixed** population parameter of interest (**which is unknown, and we aim to estimate!**).

To exemplify this, I will retake **Exercise 4** from [DSCI 552](#). For the sake of this review, suppose we want to make inference on the population tree diameter in the **SUNSET** neighbourhood in Vancouver. The dataset **vancouver_trees**, from package **datateachr**, has all population data for this neighbourhood. We can quickly obtain the population tree diameter in **SUNSET** which is 27.38 cm. Again, for the sake of this exercise, let us assume we do not know this population diameter in practice. Hence, we draw $m = 200$ random samples from this population (each with size $n = 400$) and compute the corresponding bootstrap 90%-CIs. The below table and plot indicate the following:


- The plot shows the real population diameter of 27.38 cm as a blue vertical dashed line. The $m = 200$ 90%-CIs (from random samples of size $n = 400$ each) are plotted as horizontal lines. I highlighted in red those CIs that do not contain the true population mean diameter.
- Using the table, we can see that 90.5% of this $m = 200$ CIs actually captured the true population mean diameter. **That is the concept of the frequentist CI in action!**

► [Show code cell source](#)

A tibble: 2 × 3

captured	samples	proportion
<chr>	<int>	<dbl>
FALSE	19	0.095
TRUE	181	0.905

90% Bootstrap Confidence Intervals on Tree Diameter in Sunset Neigh





Exercise 2

A CI Example with Heights

Suppose I measure the height X_i ($i = 1, 2, \dots, n$) of a **simple randomly selected** (i.e., all subjects have the same probability of being selected) subset of students in this room. My **population of interest** is the current MDS cohort.

Assume I know our true population standard deviation σ , and I want to estimate our unknown mean μ with the sample mean

$$\bar{X} = \frac{1}{n} \sum_{i=1}^n X_i.$$

I design my 95% confidence interval using the Central Limit Theorem (CLT) formula:

$$\bar{X} - 1.96 \frac{\sigma}{\sqrt{n}} \leq \mu \leq \bar{X} + 1.96 \frac{\sigma}{\sqrt{n}},$$

where 1.96 is the 0.975-quantile of the Standard Normal distribution.

Is this a valid confidence interval for μ ?

A. Yes.

B. No.



Solution to [Exercise 2](#)

No! This is an **asymptotic** CI assuming we have enough data that the sample mean is roughly normal by CLT.

Moreover, our height distribution is probably not unimodal!

Now, let us proceed with an open-ended question.



Exercise 3

A Financial CI Example



Say you collect some yearly **log returns** X_1, \dots, X_n for a financial asset, and use $\bar{X} = \frac{1}{n} \sum_{i=1}^n X_i$ to estimate the population mean.

But real **log returns** exhibit “very large/small values” much more frequently than a Normal distribution would predict. Your boss knows this and says:

“I know what we’ll do! We’ll model each X_i with a Cauchy distribution.”

Therefore, you **might** (well...) be able to compute a more suitable CI.

Nevertheless, you realize that [Cauchy distributions](#) have an undefined variance!

How do you compute a confidence interval?



Solution to [Exercise 3](#)

Here, we cannot even use our usual fallback of asymptotic intervals. We could use bootstrapping, but we can do better than that with an alternative inferential tool!

! Important

As a final thought now, before getting into our Bayesian introduction:


Even a minor change to a simple problem makes things significantly harder for frequentist methods!!!!

2. Bayesian Statistics

The **Bayesian statistical approach** addresses these challenges.

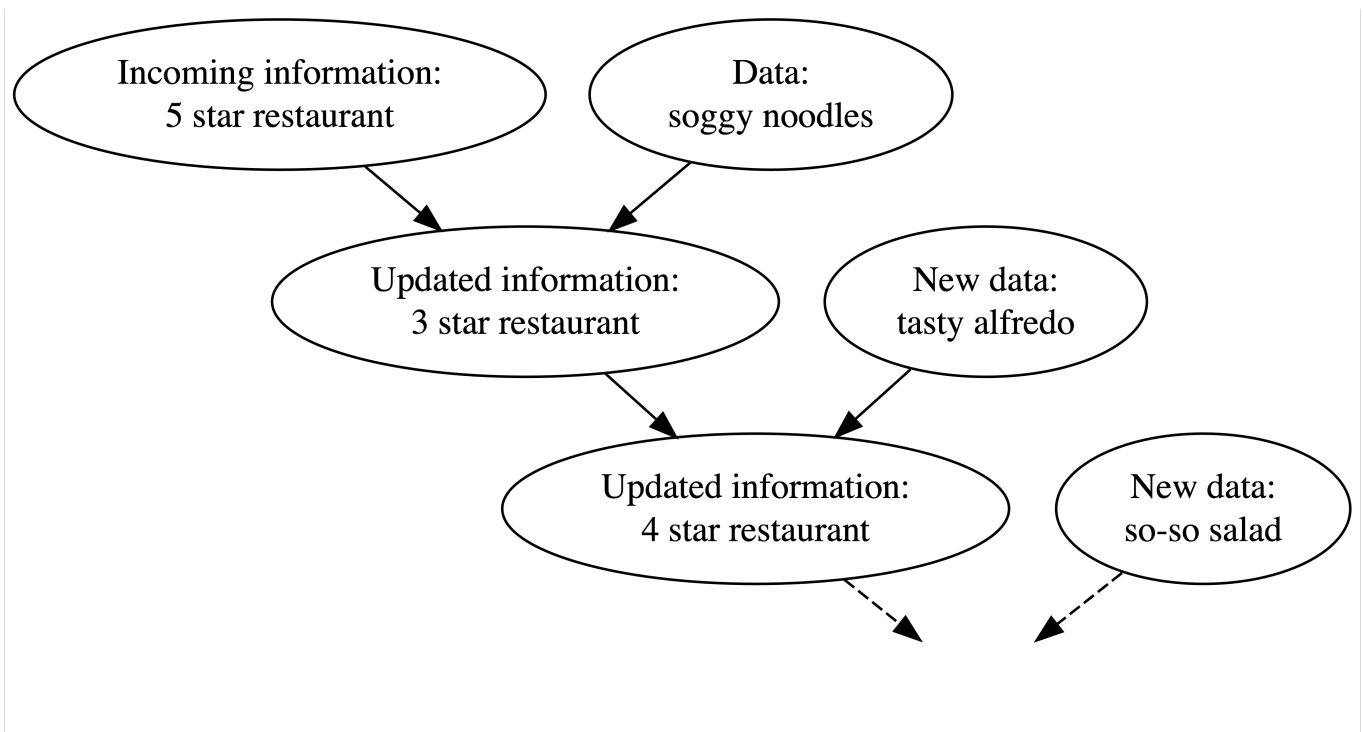
- Very flexible, handles most analysis settings.
 - Missing data, non-standard data types, non-iid, weird loss functions, adding expert knowledge.
 - *Every problem is the same, but your computer might disagree...*
- Valid inference for any (finite) amount of data.
- **Now the population parameters of interest are random variables!** They have prior and posterior distributions.
- Easy to interpret uncertainty for the population parameters.
- Posterior distribution is a “one-stop-shop” for **prediction, inference, decision-making**, etc.
- Recursive updating.

What is recursive updating?



Source: [Johnson et al. \(2021\)](#)

Let us put this concept with an example of an Italian restaurant:



Source: [Johnson et al. \(2021\)](#)

! Attention


Nonetheless, there is no free lunch...

- We would often need to be much more careful about computation.
- It requires slightly more model design work (prior).
- Less widely used in some areas of practice, e.g., medicine.
 - This is changing over time!

3. Generative Models

- **For both frequentist and Bayesian statistics**, a generative model is a simplified mathematical model for some reality: a “story about how your data were created.”
- The term **generative** means the model can be used to make *synthetic* data (i.e., running computational simulations!)

3.1. First Example



Newton's laws of dynamics $F = m \times a$ can be used to simulate position versus time given forces.

- **What if I had noisy measurements of F , m , and a and wanted to infer their true values?**

3.2. Second Example

We can model $f(\cdot)$ for whether it will rain today given a vector of **conditions/features** \mathbf{X} , $f(\mathbf{X}) \rightarrow \{\text{yes, no}\}$. What class of response is $f(\mathbf{X})$?

$f(\mathbf{X})$ is Bernoulli type!

Nonetheless, we cannot account for **everything**. Maybe only temperature, pressure, cloud cover, etc. Then, how do we deal with uncertainty in rain even under same measured conditions?

We can incorporate randomness into this model! This is analogous to ϵ on the right-hand side of a Ordinary Least-Squares (OLS) regression equation.

3.3. Probabilistic Generative Models

These models incorporate randomness into the system of interest. We design them using **probability theory** to add **wiggle room** to everything:

- We can incorporate **noise in measurements** (e.g., outputs coming from the $F = m \times a$ model).
- They can be **overly simplified models with incomplete measurements** (e.g., rainy day model).
- They can even incorporate **unobservable latent variables** (e.g., hypothetical tennis rankings).

3.4. A Probabilistic Generative Model with a Bottle Cap Flip

You have arrived early to a movie with a friend, and have great seats. However, both of you need to use the washroom. Initially, you decide to flip a coin to see who gets to go first and who will watch the seats.

But it is 2025 and nobody carries coins any more! All you have is a bottle cap from your drink:



Let us start with a FREQUENTIST inferential problem:

What is the probability π the bottle cap will land right side up? We can use n trial tosses as data.

! Attention

Let us reflect about the implications of using n trial tosses as data. We will obtain an estimate $\hat{\pi}$ as a relative **frequency** coming from this data. Hence, the name **frequentist**.

The Math of Generative Model for a Bottle Cap Flip

Any probabilistic generative model involves setting up your **random variable** of interest along with your assumed distribution.

Thus, we have the following:

$$X_i \sim \text{Bernoulli}(\pi) \quad \text{for } i = 1, \dots, n.$$

- $\pi \in [0, 1]$ is the **unknown parameter** we want to estimate: **the probability the bottle cap will land right side up**.
- X_i are the results of each bottle cap toss (1 = right side up or 0 = upside down).
- $X_i = 1$ is the success with probability π .

The **Bernoulli approach** is a **basic model** for this case; we could design another more complex probabilistic model that considers initial conditions, wind, and collisions!

4. Stan and rstan Basics

The bottle cap example was set up as a **FREQUENTIST** inferential problem (*an unknown **FIXED** parameter π we aim to estimate*). This probabilistic generative model is a Bernoulli trial.

Since the frequentist approach relies on repeating this bottle cap toss many times to estimate π , i.e., $n \rightarrow \infty$ (**not feasible to do it during our lecture time!**), we will computationally simulate this system with a Monte Carlo simulation in a worksheet (not for marks) during

lab1.

By definition, a Monte Carlo simulation will computationally repeat the event of interest n times with random inputs to obtain our n outputs of interest (a binary outcome in this case).



We can easily do this via the base `R` function `rbernoulli()` from `purrr`. But we will use this example to introduce `Stan` and `rstan`.

`Stan` is a probabilistic `C++`-based programming language for Bayesian statistical inference (but it can also perform simple Monte Carlo simulations). The `R` package `rstan` will allow pulling the simulation outputs from `Stan`.

4.1. Coding the Model and Running your Simulation

In general, we will follow these steps:

1. Code up our generative model in `Stan`.
2. Specify observed values of data to estimate using `rstan` (not necessary for Monte Carlo simulations).
3. Generate **synthetic data**.
4. Perform inference with your simulation outputs.

The generative model is all you need (and all you get!)

- Once you have a generative model, you can derive **everything**: tests, inference, etc.
- If your model **can** generate it, it will be handled in inference:
 - missing data, dependence, complex data types, etc.
- If your model **cannot** generate it, it **will not be handled**.

4.2. Can We Theoretically Estimate π ?

Yes, we can! **Maximum likelihood estimation** (MLE) is the standard frequentist approach. Recall the procedure from [DSCI 551's lecture7](#).

How do we find the maximum likelihood estimate for π between 0 and 1?

- We set up the Bernoulli trial for one toss:

$$\ell(\pi \mid x_i) = P(X_i = x_i \mid \pi) = \pi^{x_i}(1 - \pi)^{1-x_i} \quad \text{for } x_i = 0, 1.$$

- We obtain the joint likelihood function with n iid tosses:

$$\begin{aligned} \ell(\pi \mid x_1, \dots, x_n) &= \prod_{i=1}^n \pi^{x_i}(1 - \pi)^{1-x_i} \\ &= \pi^{\sum_{i=1}^n x_i} (1 - \pi)^{n - \sum_{i=1}^n x_i}. \end{aligned}$$

- Then, the log-likelihood function:

$$\log \ell(\pi \mid x_1, \dots, x_n) = \sum_{i=1}^n x_i \log(\pi) + \left(n - \sum_{i=1}^n x_i \right) \log(1 - \pi).$$

- We take the first partial derivative of the log-likelihood function with respect to π , set it to 0, and solve for π :

$$\hat{\pi} = \frac{1}{n} \sum_{i=1}^n X_i.$$

4.3. Uncertainty in π

How can we characterize the uncertainty in π (without doing boatloads of math)?

A bootstrap CI might be the answer, but what if we only have $n = 10$ replicates?

```
options(repr.matrix.max.rows = 10)
library(purrr)
library(tidyverse)
library(infer)

# Obtaining sample of size n = 10
set.seed(553)
sample_n10 <- tibble(flip = as.character(rbernoulli(n = 10, p = 0.7)))
sample_n10
```

Warning message:
“`rbernoulli()` was deprecated in purrr 1.0.0.”

A tibble:

10 × 1

flip
<chr>
TRUE
TRUE
TRUE
TRUE
TRUE
TRUE
TRUE
TRUE
FALSE
TRUE

```
# Obtaining bootstrap distribution via infer
set.seed(553)
bootstrap_distribution <- sample_n10 |>
  specify(response = flip, success = "TRUE") |>
  generate(reps = 1000, type = "bootstrap") |>
  calculate(stat = "prop")
mean_stat <- mean(bootstrap_distribution$stat)
mean_stat

# Obtainig percentile 95% CI
percentile_CI <- bootstrap_distribution |>
  get_confidence_interval(level = 0.95, type = "percentile")
percentile_CI
```

0.8967

A tibble: 1 × 2


lower_ci	upper_ci
<dbl>	<dbl>
0.7	1

```
options(repr.plot.height = 7, repr.plot.width = 12)

bootstrap_plot <- bootstrap_distribution |>
  visualize(bins = 6) +
  shade_confidence_interval(endpoints = percentile_CI) +
  geom_vline(xintercept = mean_stat, linetype = "dashed") +
  theme(
    plot.title = element_text(size = 24, face = "bold"),
    axis.text = element_text(size = 17),
    axis.title = element_text(size = 21)
  ) +
  labs(x = "Statistic", y = "Count")

bootstrap_plot
```

Simulation-Based Bootstrap Distribution



! Important

Our 95% bootstrap CI shows **biased results**; the lower bound is 0.7! This is not good from an inferential perspective. An easy solution would be increasing the sample size n , but what if this is not possible?

Bayesian inference will open up our set of possible solutions to this matter.

5. Difference between Probability and Likelihood

It is important to emphasize that in Statistics, **probability** and **likelihood** are **NOT** the same. In general, **probability** refers to the chance that some outcome of interest will happen for a particular random variable. Note a probability is always bounded between 0 and 1. Conversely, **given some observed data**, a **likelihood** refers to how **plausible** a given **distributional parameter** is. Furthermore, a likelihood is not bounded between 0 and 1.

Let us explore this concept via the Binomial distribution. Let X be the number of successes after n independent Bernoulli trials with probability of success $0 \leq \pi \leq 1$. Then, X is said to have a Binomial distribution:

$$X \sim \text{Binomial}(n, \pi).$$

A Binomial distribution is characterized by the probability mass function (PMF)

$$P(X = x \mid n, \pi) = \binom{n}{x} \pi^x (1 - \pi)^{n-x} \quad \text{for } x = 0, 1, \dots, n. \quad (1)$$



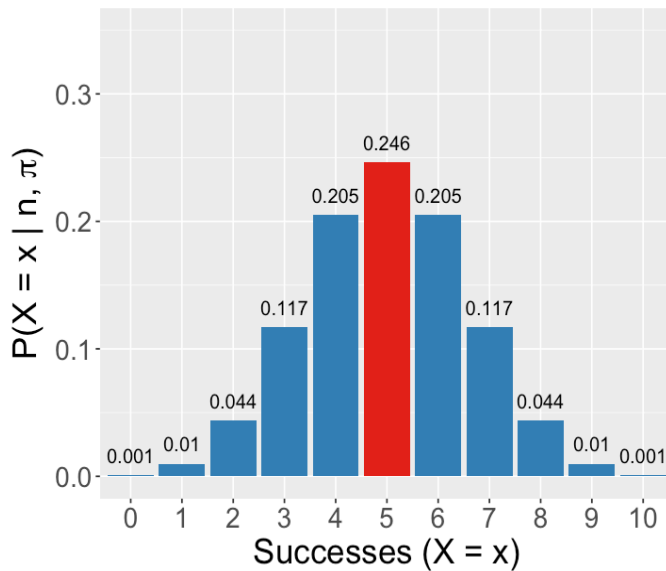



Term $\binom{n}{x}$ indicates the total number of combinations for x successes out of n trials:

$$\binom{n}{x} = \frac{n!}{x!(n-x)!}.$$

Let us plot the PMFs of six Binomial random variables. Note these plots indicate a **probability** $P(X = x \mid n, \pi)$ on their y -axes and they have the same number of trials $n = 10$:

- $n = 10$ and $\pi = 0.3$.
- $n = 10$ and $\pi = 0.4$.
- $n = 10$ and $\pi = 0.5$.
- $n = 10$ and $\pi = 0.6$.
- $n = 10$ and $\pi = 0.7$.
- $n = 10$ and $\pi = 0.8$.

Suppose we are specifically interested in the outcome $Y = 5$, highlighted as a red bar in the six PMFs. We also indicate the **probabilities** associated to each outcome on top of each bar.

PMF of a Binomial($n = 10, \pi = 0.3$)PMF of a Binomial($n = 10, \pi = 0.4$)PMF of a Binomial($n = 10, \pi = 0.5$)PMF of a Binomial($n = 10, \pi = 0.6$)PMF of a Binomial($n = 10, \pi = 0.7$)PMF of a Binomial($n = 10, \pi = 0.8$)

Suppose we sample data from a given Binomial population of interest, and this sample is composed of $n = 10$ trials and $x = 5$ successes. Under a frequentist paradigm, we are interested in inferring that value of π , which is the **most likely** for these values of n and x . Then, this is what we can do:

- Our PMF (1) will become a likelihood function **given** $n = 10$ and **our observed** $x = 5$ successes (note the lowercase x):


$$\begin{aligned}\ell(\pi \mid x = 5, n = 10) &= P(X = 5 \mid n = 10, \pi) \\ &= \binom{10}{5} \pi^5 (1 - \pi)^{10-5}.\end{aligned}\tag{2}$$

! Attention

The likelihood (2) is **mathematically** equal to the Binomial PMF with $n = 10$ and $X = 5$. Nevertheless, this likelihood is now **in function** of the parameter π .

- That said, the below plot shows the likelihood function (2) with the plausible range of π on the x -axis (which is bounded between 0 and 1 since it is the probability of success in the Binomial distribution). Moreover, the likelihood values are indicated in the y -axis.

Binomial likelihood function of π given observed $x = 5$ successes and $n = 10$ trials



In the above likelihood plot, we highlighted as red vertical dashed lines those values corresponding to the red bars in the previous six Binomial PMFs (**since a probability is MATHEMATICALLY equal to a likelihood, but not STATISTICALLY**).

The plot shows that the likelihood function [\(2\)](#) gets maximized when $\pi = 0.5$. That said, given our observed collected sample with $x = 5$ successes with $n = 10$, a value of $\pi = 0.5$ is **the most plausible (or likely!)** in this Binomial population.

6. Wrapping Up with Frequentist Drawbacks

- We cannot incorporate our “expert” intuition, probabilistically speaking.
- How do we use side information, e.g., results of flips from other bottle caps?
- Asymptotic tools tend not to work well with small amounts of data.